

WILFREDO PAULO A. PEREZ III

Naga City, Camarines Sur, Philippines 4400

☎ +63 956-943-6779 | ✉ pauloperez9754@gmail.com | 🔗 linkedin.com/in/wppereziii | 🌐 github.com/C0deRhin0

EDUCATION

Ateneo de Naga University

2022 – 2026

Candidate for Magna Cum Laude - Bachelor of Science in Computer Science

Naga City, Camarines Sur

- **Thesis:** A Vendor-Agnostic, Resource-Aware, Explainable Software-Defined Intrusion Prevention System for SOHO Networks using PPO, Policy Distillation, and TreeSHAP-IQ
- Current QPI: 3.64; Consistent President's Lister; DOST Scholar

WORK EXPERIENCE

Mindrift

November 2025 – April 2026

AI Engineer (Task-Based, Flexible) — LLM Training & Data

Remote

- Contribute to large language model (LLM) training pipelines by sourcing, scraping, and curating structured web data; have delivered **3,200+ annotated data samples** across multiple task batches for fine-tuning datasets.
- Engineered a **fully automated web scraping pipeline** (Python, BeautifulSoup, Selenium) that reduced per-task manual processing time by **~85%**, enabling consistent high-volume output without interrupting primary employment duties.
- Maintained a **94% task acceptance rate** across submitted batches, reflecting strong data quality standards aligned with LLM training requirements.
- Operates independently across flexible hours, demonstrating capacity to manage **concurrent professional obligations** with zero scheduling conflicts.

Nueca Technologies Inc.

September 2025 – April 2026

Cybersecurity Associate (Part-Time) — Cybersecurity Department

Onsite

- Administer the company's internal LAN serving **55 employees across 5 departments**, enforcing segmentation, access control, and 99%+ uptime of critical network services.
- Manage an additive email security layer that has **intercepted and blocked 100+ malicious emails** (phishing, spoofing) since deployment, reducing end-user phishing exposure by an estimated **70%**.
- Designed and delivered **3 company-wide cybersecurity orientation sessions** covering 100% of staff (55 employees), reinforcing CIA triad principles and reducing self-reported security incident frequency.
- Appointed **Compliance Officer for Privacy (COP)** and led the company's NPC Seal of Registration submission, coordinating data privacy compliance across all 5 departments.
- Provided compensated security consulting to **2 partner organizations**, identifying and remediating critical vulnerabilities and delivering formal risk reports.
- Architected an **AI orchestration platform** targeting a projected **30% reduction in repetitive operational overhead** across Engineering and Operations workflows.
- Led the development of an internal **ERP system with integrated IAM, HR management, and asset tracking modules**, designed to consolidate workflows for all 55 users and reduce manual administrative overhead by an estimated **60%**.
- Proposed and currently implementing **4 network security enhancements**, including NAC policy hardening and traffic monitoring improvements, as part of the department's ongoing infrastructure roadmap.

Nueca Technologies Inc.

May 2025 – July 2025

Network and Cloud Security Analyst [Intern] — Cybersecurity Department

Onsite

- Conducted a full network audit of legacy infrastructure; identified **12 segmentation and visibility gaps** and proposed a modular Layer 3 redesign with VLAN isolation, centralized firewalls, and real-time monitoring – **approved and funded for implementation**.
- Performed vulnerability scanning across **3 network segments** using Nmap and OpenVAS; delivered a prioritized remediation plan covering **20+ identified risks**, ranked by CVSS severity.
- Audited **15 macOS workstations** against CIS Benchmarks using automated and manual checks; achieved **82% compliance rate** post-remediation (up from 51%) within a 2-week cycle.
- Deployed Wazuh SIEM and HIDS in a proof-of-concept lab (Manager, Indexer, agents), enabling centralized logging and rule-based alerting across **5 monitored endpoints**.
- Built secure Dockerfiles and container stacks with Docker Compose; migrated **3 services** to Kubernetes using declarative manifests (Deployments, Services, RBAC, ConfigMaps, Secrets, probes, HPA, Ingress).
- Implemented production-grade Wazuh SIEM on managed Kubernetes (DigitalOcean) using StatefulSets, DaemonSets, TLS, and Kibana visualization, supporting **real-time log aggregation across 8 nodes**.
- Resolved **3 critical DNS misconfigurations** (missing SPF, DKIM, DMARC records), directly eliminating the company's domain spoofing vulnerability.
- Delivered a live macOS spyware attack simulation at internal DevCon for **20+ engineering staff**, demonstrating endpoint compromise vectors and contributing to ISO-aligned policy development.

PROJECTS

CheatScale: Enterprise AI Development Orchestration | *OpenCode, Python, Obsidian*

February 2026 – May 2026

- Built an enterprise-grade AI development orchestration system featuring **24 specialized AI agents** operating across a 3-tier architecture (Presentation, Logic, Data), enabling parallel frontend/backend agent execution with clean separation of concerns.
- Integrates wave-based task routing, Obsidian graph-linked documentation across **120+ project nodes**, and automated roadmap generation — reducing project scaffolding overhead by an estimated **88%**.
- Enforces **80%+ unit test coverage** requirements by isolating domain logic from rendering layers, making AI-generated code independently testable and production-auditable.

NuecaAI RAG Customer Support Chatbot | *n8n, Qdrant, OpenAI Embeddings, Node.js, Python*

April 2026 – May 2026

- Built and deployed a **production RAG-powered customer support chatbot** for Nueca Technologies using n8n workflow automation, Qdrant vector database, and OpenAI embeddings for semantic FAQ retrieval — actively serving company operations across a **structured multi-turn web interface**.
- Engineered a Python FAQ ingestion pipeline and a Node.js proxy server handling real-time webhook sessions; reduced average support query resolution time by an estimated **~70%** by replacing manual FAQ lookup with **sub-second vector similarity search** across the company's full knowledge base.

NuecaAI Receipt Hybrid OCR | *React, Vite, Node.js, Claude Vision, Tesseract.js, Ollama, Tailwind CSS*

April 2026 – May 2026

- Engineered a dual-mode receipt digitization system for Nueca Technologies with intelligent cloud routing (Claude Vision: **~95%+ accuracy, 3–5s per receipt**) and a fully local privacy fallback (Tesseract.js + Ollama llama3.2:3b: **~75% accuracy, 5–15s**), deployed as an **HTTPS LAN server** accessible across all company devices without external configuration.
- Reduced manual receipt data-entry time by an estimated **~92% per document**; outputs structured JSON (vendor, date, line items, totals) with **one-click CSV and Google Sheets export**, and eliminates **100% of cloud data exposure** in secure mode for privacy-sensitive financial processing.

NuecaAI Whisper Local — Offline Meeting Transcriber | *whisper.cpp, Ollama, pyannote.audio, React, Python*

April 2026 – May 2026

- Built a **100% offline** meeting transcription and record-extraction system for Nueca Technologies using whisper.cpp with Metal GPU acceleration (Apple Silicon); implemented **2 parallel transcription workers** achieving a **~2× throughput improvement** over single-threaded baselines on multi-core machines.
- Handles **60+ minute recordings** via smart silence-based chunking; integrates pyannote.audio for neural speaker diarization with structured, speaker-labeled output exportable as **styled PDFs with page numbers** — eliminating 100% of cloud data exposure for sensitive internal meetings and reducing post-meeting documentation time by an estimated **~75%**.

Corp-Mind-AI — On-Premise HR Knowledge Assistant | *FastAPI, Qdrant, Ollama, React, Vite, sentence-transformers*

March 2026 – May 2026

- Architected a **fully on-premise RAG HR assistant** for Nueca Technologies using hybrid semantic + BM25 search (Qdrant RRF fusion), delivering source-cited answers with **page-level document traceability** across the company's internal corpus — zero bytes of employee data leave the network.
- Implemented a **2-hour TTL answer cache** reducing redundant LLM inference overhead by an estimated **~65%** for repeated HR queries; features an **admin panel** for web-based document ingestion and a diagnostics dashboard monitoring Qdrant, Ollama, and collection health in real time.

Vector-Mind-AI — Multi-Agent Research Orchestrator | *Ollama, Qdrant, FastAPI, React, OpenAI, Anthropic API*

March 2026 – Present

- Built a **6-agent wave-based research orchestration platform** (Planner → Researcher → Analyst → Writer → Critic → Fact Checker) with cross-session knowledge graph visualization, supporting interchangeable local (Ollama) and cloud (OpenAI, Anthropic, Google) LLM backends — an open-source, local-first alternative to Google NotebookLM.
- Generates polished, multi-format outputs (research briefs, blog posts, executive summaries, presentation scripts, Twitter threads) from a single query; integrates RAG with persistent knowledge graph linking, reducing manual research synthesis time by an estimated **~65%** versus traditional workflows across **6+ output format templates**.

LED-Entropy: Hardware True Random Number Generator | *Raspberry Pi, Python, Electronics*

October 2025 – April 2026

- Engineered a hardware entropy source combining analog capacitor/transistor noise, photoresistor sensing, and chaotic LED patterns captured via Raspberry Pi camera, a miniaturized analogue of Cloudflare's LavaLamp system.
- Generates **2,400+ cryptographically-useful random bits/second** from physical phenomena, achieving **~38% lower entropy bias** than software PRNGs under equivalent benchmark conditions.

AI-Centric Email Security Responder | *Power Automate, n8n, LLM API*

October 2025 – April 2026

- Automated end-to-end email security response drafting using Power Automate and n8n orchestration, integrating an LLM API to generate context-aware ESL responses.
- Reduced average security response drafting time by **~76%** across **3 monitored mailboxes**, eliminating an estimated **5 manual hours per week** of analyst triage and reply work.

Home-Laboratory Software-defined Firewall | *pfSense, Suricata, OSI Layer 1–4*

April 2025 – June 2025

- Deployed a software-defined firewall across a **4-VLAN segmented home-lab**, enforcing **60+ custom firewall rules** with policy-based traffic control and real-time IDS visibility via Suricata.
- Achieved **zero lateral movement** between segments during internal threat simulations; reduced unfiltered east-west traffic by **~91%** versus the flat-network baseline.

AWS Cloud DevSecOps | *AWS, GitHub Actions, IaC*

March 2025 – May 2025

- Provisioned a fully automated CI/CD pipeline across **7 AWS services** (S3, EC2, IAM, CodeArtifact, CodeBuild, CodeDeploy, CodePipeline) via IaC, reducing manual deployment steps from **14 to 2**.
- Cut average build-to-deploy cycle time by **~72%** versus manual baseline; enforced least-privilege IAM policies across **4 defined roles** with zero standing admin access.

AdNU CEVAS Certificate System | *LAMP, Docker, Kubernetes*

November 2024 – January 2025

- Built a LAMP-stack certificate validation platform for the Ateneo de Naga CEVAS institute, supporting **3 academic departments** and the validation of **200+ certificates** using SHA-256 cryptographic hashing.
- Reduced certificate authenticity verification time from **~3 days (manual process)** to **under 10 seconds**; containerized for portable deployment via Docker or Kubernetes.

Offensive Python Scripts | *Python, MITRE ATT&CK*

March 2025

- Developed **11 MITRE ATT&CK-mapped attack simulations** spanning **5 tactics** (Reconnaissance, Execution, Persistence, Credential Access, Exfiltration), each paired with detection logic and mitigation guidance for red-blue team training use.

Automated System Security Logging | *Bash, PowerShell*

March 2025

- Developed cross-platform forensic logging scripts for Linux and Windows, aggregating **18 log source categories** (auth, process, network, file integrity) into structured reports in under **90 seconds per host**.
- Reduced manual log collection and triage time by **~80%** versus ad-hoc shell commands; outputs are directly ingestible by Wazuh and Splunk.

ACTIVITIES

Certifications: ISC² Certified in Cybersecurity (CC); CCSP-AWS; CNSP; BTL0; FIT-SM — *In Progress: Google Cybersecurity Professional Certificate, CompTIA Security+*

Competitions: Regional Hack4Gov 2 & 3 — AdNU Representative; National Hack4Gov 3 — AdNU Representative

ADDITIONALS

Programming Languages: C, C++, Python, Java, React, TypeScript, JavaScript, Ruby

Languages: English (B2: Upper-Intermediate), Tagalog (Native)

Technical Skills: AI Harness, Docker, Kubernetes, AWS, DigitalOcean, Wazuh SIEM, Nmap, OpenVAS, Suricata, Wireshark, TCPdump, pfSense, MITRE ATT&CK, Nessus, Linux Shells, AWK/GAWK, Maltego, Google Dorking, Solidity, n8n, Qdrant, Ollama, Whisper, FastAPI

Soft Skills: Data-driven strategic planning, Organized, Time Management, Professional, Cooperative

References: **available upon request*